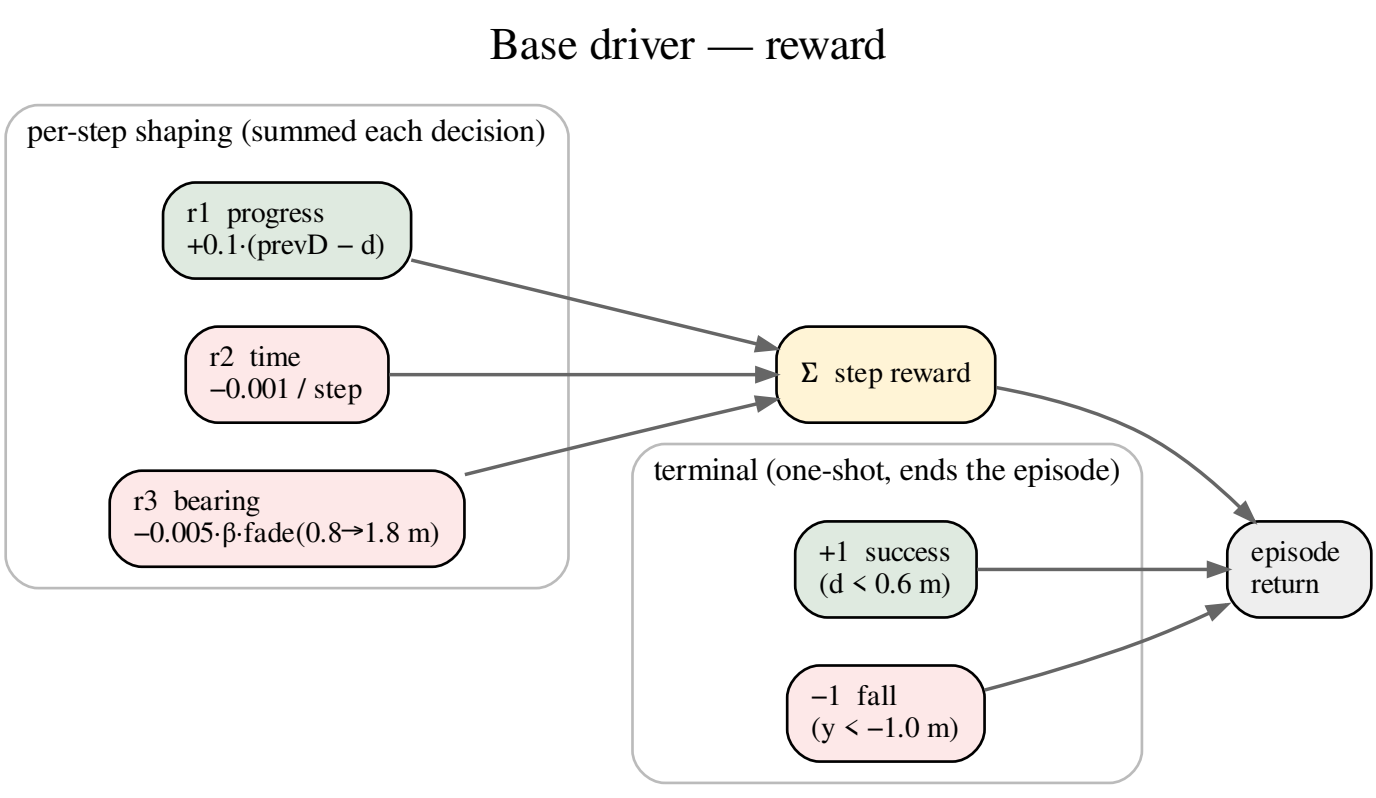
Base driver — reward
per-step shaping (summed each decision)
r1 progress
+0.1·(prevD − d)
r2 time
−0.001 / step
r3 bearing
−0.005·β·fade(0.8→1.8 m)
Σ step reward
terminal (one-shot, ends the episode)
+1 success
(d < 0.6 m)
−1 fall
(y < −1.0 m)
episode
return